
Protocol Design Space of Snooping Cache Coherent Multiprocessors

Recap

- **Snooping cache coherence**
 - solve difficult problem by applying extra interpretation to naturally occurring events
 - » state transitions, bus transactions
 - write-thru cache
 - » 2-state: invalid, valid
 - » no new transaction, no new wires
 - » coherence mechanism provides consistency, since all writes in bus order
 - » poor performance
- **Coherent memory system**
- **Sequential Consistency**

Sequential Consistency

- **Memory operations from a proc become visible (to itself and others) in program order**
- **There exist a total order, consistent with this partial order - i.e., an interleaving**
 - the position at which a write occurs in the hypothetical total order should be the same with respect to all processors
- **Sufficient Conditions**
 - every process issues mem operations in program order
 - after a write operation is issued, the issuing process waits for the write to complete before issuing next memory operation
 - after a read is issued, the issuing process waits for the read to complete and for the write whose value is being returned to complete (globally) before issuing its next operation

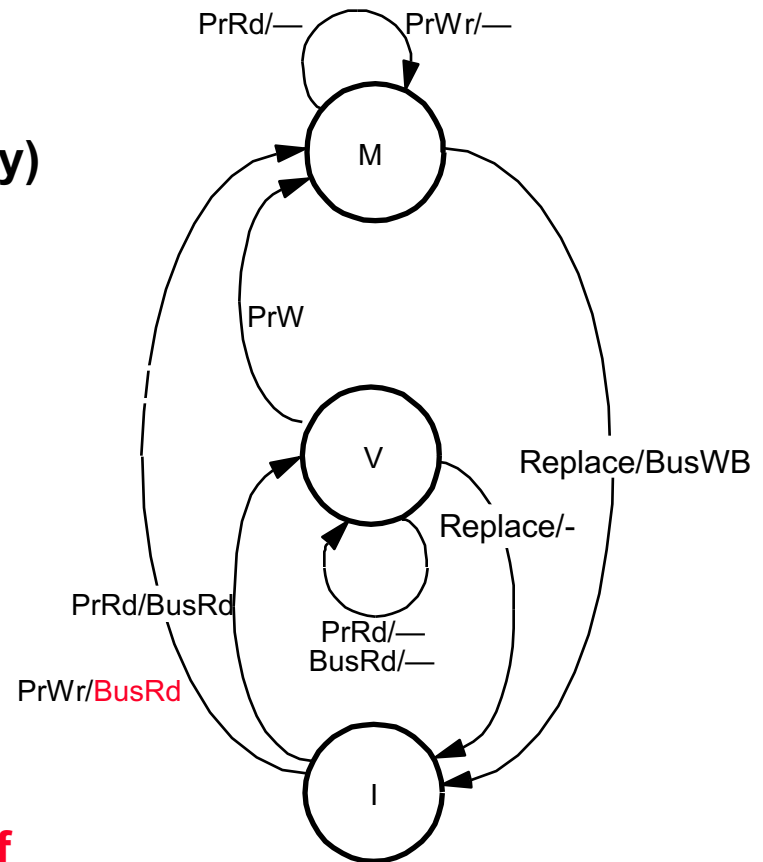
Write-back Caches

- **2 processor operations**
 - PrRd, PrWr
- **3 states**
 - invalid, valid (clean), modified (dirty)
 - ownership: who supplies block
- **2 bus transactions:**
 - read (BusRd), write-back (BusWB)
 - only cache-block transfers

=> treat Valid as “shared” and Modified as “exclusive”

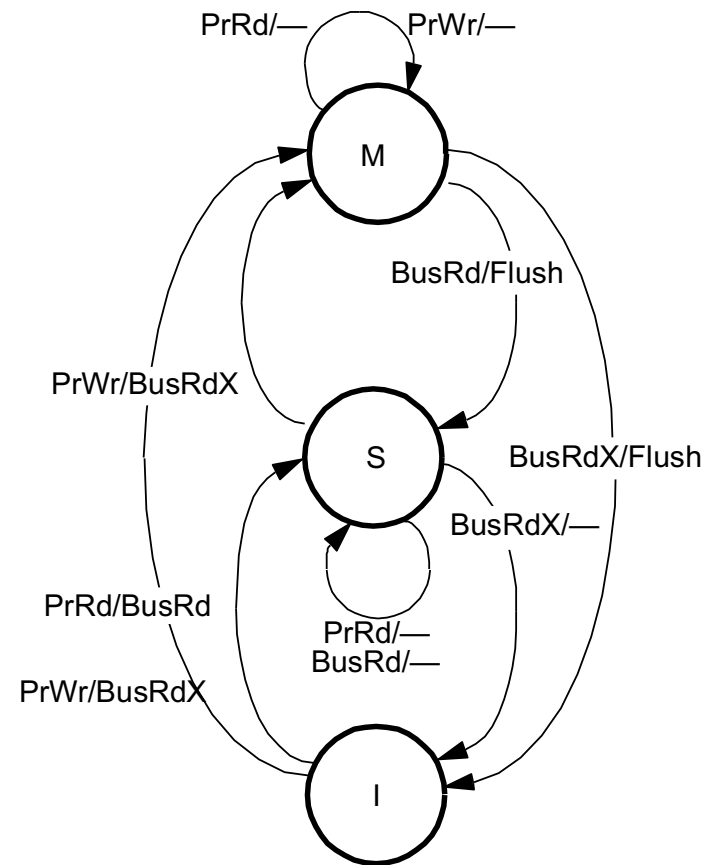
=> introduce one new bus transaction

- **read-exclusive: read for purpose of modifying (read-to-own)**

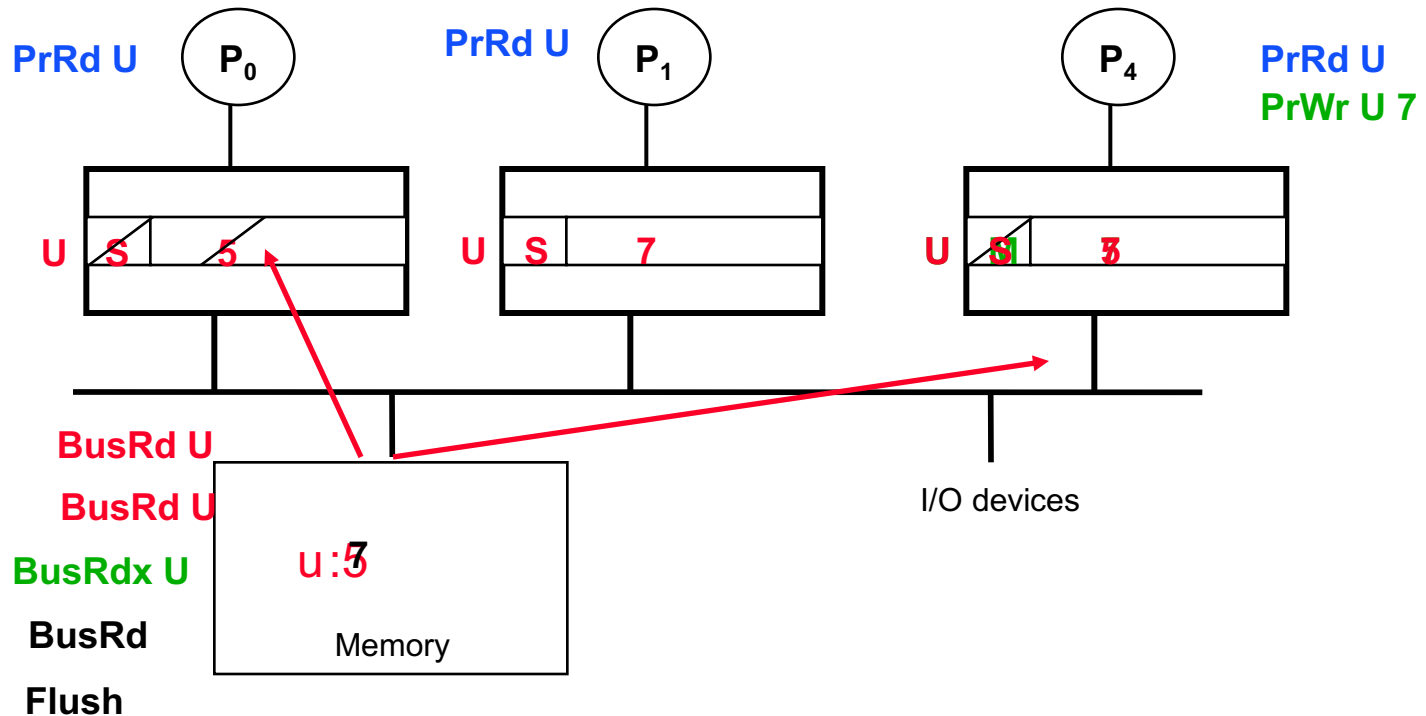


MSI Invalidate Protocol

- **Read obtains block in “shared”**
 - even if only cache copy
- **Obtain exclusive ownership before writing**
 - BusRdx causes others to invalidate (demote)
 - If M in another cache, will flush
 - BusRdx even if hit in S
 - » promote to M (upgrade)
- **What about replacement?**
 - S->I, M->I as before



Example: Write-Back Protocol



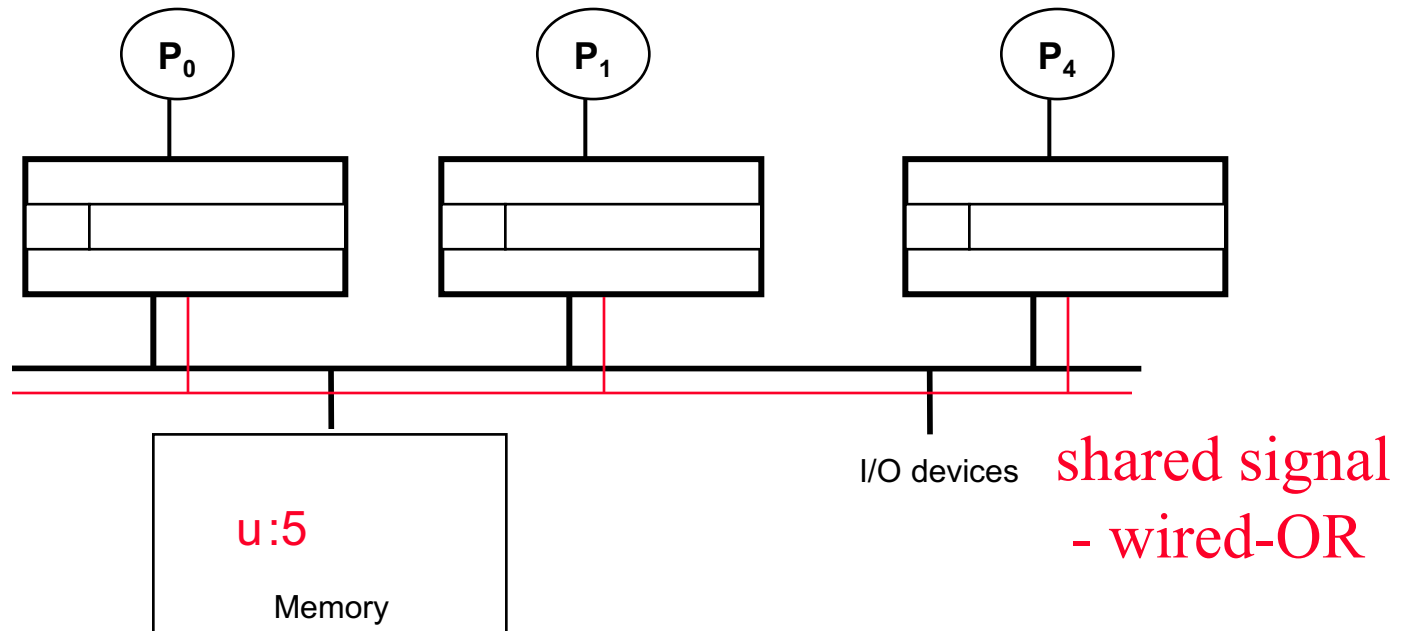
Lower-level Protocol Choices

- **BusRd observed in M state: what transition to make?**
 - M ----> I
 - M ----> S
 - Depends on expectations of access patterns
- **How does memory know whether or not to supply data on BusRd?**
- **Problem: Read/Write is 2 bus xactions, even if no sharing**
 - » BusRd (I->S) followed by BusRdX or BusUpgr (S->M)
 - » What happens on sequential programs?

MESI (4-state) Invalidation Protocol

- **Add *exclusive* state**
 - distinguish exclusive (writable) and owned (written)
 - Main memory is up to date, so cache not necessarily owner
 - can be written locally
- **States**
 - invalid
 - exclusive or *exclusive-clean* (only this cache has copy, but not modified)
 - shared (two or more caches may have copies)
 - modified (dirty)
- **I -> E on PrRd if no cache has copy**
=> How can you tell?

Hardware Support for MESI



- All cache controllers snoop on BusRd
- Assert 'shared' if present (S? E? M?)
- Issuer chooses between S and E
 - how does it know when all have voted?

MESI State Transition Diagram

- BusRd(**S**) means shared line asserted on BusRd transaction
- Flush': if cache-to-cache xfers
 - only one cache flushes data
- MOESI protocol: Owned state: exclusive but memory not valid

